

KurveRSC

Validation-Guided Relational Signal Compression with a Downstream Learner in the Loop

Kurve AI

September 1, 2026

Contents

1	Abstract	1
2	Thesis	2
3	Difference from relational feature synthesis	2
4	Introduction	3
5	Relational signal compression	4
6	GraphReduce execution algorithm	6
7	Relational feature construction	8
8	KurveRSC meta-optimization	10
9	Final fitting and prediction	13
10	Systems implementation	15
11	Experimental protocol and results	15
12	Related work	22
13	Limitations and future work	22
14	Conclusion	23
15	References	23
16	Appendix A: Reference API	24
17	Appendix B: Default meta-parameter frontier	25
18	Appendix C: Artifact and trial record	25

1 Abstract

Relational prediction begins with a representation problem: a database contains many tables at different granularities, whereas conventional supervised learners consume one matrix at one prediction grain. A fixed flattening policy makes that representation an engineering constant. KurveRSC instead treats it as a validation-tested model choice. KurveRSC defines a finite space of relational programs over a database compute graph. A program chooses graph depth, a cumulative set of relational feature families, automatic annotation policy, and feature budgets. GraphReduce executes each program under point-in-time cutoffs and produces an entity-grain feature frame. A downstream

learner is fitted to every frame, and its held-out predictive performance supplies the objective used to select the relational program. The result is a fitted artifact containing both a frozen GraphReduce execution plan and a predictive model.

The current reference implementation uses CatBoost and selects configurations by validation AUROC for binary classification or mean absolute error (MAE) for regression. The boundary between relational materialization and estimation is a typed tabular frame, so the learner is replaceable; an experimental TabPFN v3 adapter demonstrates this separation. The contribution is not a new tree learner or a differentiable database operator. It is an auditable optimization procedure over the *relational signal-compression space*, evaluated jointly with the learner that will consume the compressed representation.

The default reproducibility profile evaluates graph configurations on the complete latest eligible training cutoff, executes configurations and cutoff folds sequentially, reranks three finalists over three complete temporal folds, and fits the production artifact on the latest eligible training cutoff. Only one relational feature frame is resident at a time. Additional production cutoffs can improve predictive performance on some tasks, but they are opt-in so that the reference configuration remains practical to reproduce under bounded memory.

This report describes the GraphReduce execution algorithm, the relational feature families, the KurveRSC meta-parameters and selection objective, temporal leakage controls, frozen-plan lifecycle, bounded-memory implementation, and the intended RelBench/RelArena evaluation protocol. The performance section contains the complete 21-task controlled-run snapshot.

2 Thesis

The central thesis of KurveRSC is that *relational signal compression is a first-class, durable optimization surface*. Predictive performance depends not only on the learner applied after data preparation, but also on which relational paths, temporal horizons, reductions, annotations, and source columns survive the conversion from a variable-grain database into the representation consumed by that learner. A fixed flattening recipe hard-codes this choice; a monolithic relational architecture hides it inside one model. KurveRSC instead exposes the representation program, measures it jointly with the learner, and freezes the selected program as an auditable artifact.

This view preserves learner pluggability. CatBoost is the current reference judge, but the selected GraphReduce frame has a typed tabular boundary that can be consumed by TabPFN-3 or another downstream learner or foundation model. The research direction is therefore not “feature engineering instead of foundation models.” It is task-conditioned relational compression *for* whatever learner best exploits the resulting signal.

Recent results from two independent systems motivate this thesis. KumoRFM-2 reports that task-specific fine-tuning raises average MRR on SALT from 0.83 for its in-context model to 0.89, compared with 0.79 for the strongest supervised baseline in that study [15]. This is not representation-program search, but it shows that even a strong pretrained relational model retains substantial value in task-specific adaptation. Prior Labs’ TabPFN-Rel constructs relational features with Deep Feature Synthesis and delegates prediction to TabPFN-3 [10]. Its RelArena performance shows that a foundation learner can remain competitive when the relational representation is built by an independent preprocessing layer. Neither result proves KurveRSC’s mechanism; together they support the broader modular hypothesis that the relational representation and the learner should be adapted jointly without permanently binding either layer to the other.

3 Difference from relational feature synthesis

The shortest distinction is: **Deep Feature Synthesis generates a feature space; KurveRSC selects a relational program by measuring how well its complete feature frame works with the downstream learner.** DFS and KurveRSC are therefore complementary ideas at different levels. DFS defines how primitives can be composed across relationships [2]. KurveRSC defines how alternative relational representation policies are proposed, evaluated, complexity-checked, selected, and frozen.

Dimension	Deep Feature Synthesis	GraphReduce	KurveRSC
Primary object	Composed feature definitions	Executable table graph and node operations	Search distribution over complete GraphReduce programs

Dimension	Deep Feature Synthesis	GraphReduce	KurveRSC
Main choices	Relationships, primitive set, synthesis depth	Graph depth, node policy, feature families, cutoffs, reductions	Depth, family combinations, annotations, feature budgets, temporal policy, and learner policy
Downstream learner	Normally fitted after synthesis; not the DFS generation objective	External to the execution engine	In the loop: validation AUROC or MAE scores every candidate frame
Task adaptation	Caller selects recipe; downstream feature selection may remove columns	Caller configures a graph program	Search, confirmation, temporal reranking, and complexity/stability selection are target-conditioned
Temporal semantics	Depend on cutoff-time configuration supplied by the caller	Point-in-time filters and label operations are native graph operations	Cutoff-safe candidates and folds are part of selection and final replay
Output	Feature definitions and a materialized table	Reduced entity-grain frame plus operation lineage	Fitted artifact containing the selected graph configuration, frozen execution plan, schema, and learner
Inference behavior	Recompute the chosen feature definitions	Re-execute the configured graph	Replay the exact learned plan with feature discovery disabled

This distinction also separates KurveRSC from simply placing an AutoML model after a very wide DFS matrix. Column-level learner regularization can decide which already-materialized columns to use, but it cannot recover relational paths, cutoff windows, feature families, or propagation depths that were never materialized. Conversely, exhaustively generating everything can be computationally prohibitive and statistically noisy. KurveRSC searches those upstream choices under explicit width and complexity controls before committing to the final full-data representation.

RDBLearn and TabPFN-Rel demonstrate a strong DFS-to-foundation-model pipeline [9,10]. KurveRSC generalizes a different boundary: the relational program is the optimized object, and the frame consumer is replaceable. A DFS engine could in principle serve as another program generator inside the same learner-guided selection abstraction; GraphReduce is the current execution substrate because it exposes cutoff-safe graph operations and a replayable plan.

4 Introduction

Operational data is usually relational. An entity such as a customer, product, driver, study, or advertisement is connected to transactions, events, messages, measurements, and other entities through foreign keys. A supervised model, however, usually expects a fixed-width matrix whose rows share one prediction grain. Constructing that matrix requires decisions about which tables to reach, how far to traverse, what history is visible, which reductions to apply, and how much feature growth to tolerate.

Automatic relational feature engineering has a long history. Deep Feature Synthesis composes aggregation and transformation primitives across relationships [2], and OneBM develops a scalable relational feature-generation framework [3]. GraphReduce takes a graph-programming approach: tables are nodes, key relationships are edges, and a granularity-directed depth-first traversal reduces child histories before joining them into their parents [1]. This preserves a declared root grain while supporting SQL execution, temporal filters, custom operations, and automated feature families.

These systems still leave a consequential question: *which relational program should be used for this prediction task?* Greater depth reaches more tables, and richer feature families preserve more behavior, but additional features can be sparse, redundant, unstable, expensive, or actively harmful to generalization. The most elaborate graph is not necessarily the most predictive graph.

KurveRSC moves relational-program choice inside the model-selection loop. For each candidate program, it materializes a point-in-time feature frame, fits the same downstream learner policy, and evaluates the complete representation-learner pair. This yields three practical properties:

1. **Task-conditioned relational representation.** The chosen view depends on the target, split, and learner rather than on an a priori belief that more joins are always better.
2. **Auditable selection.** Every candidate records its configuration, feature count, row count, materialization time, model time, validation score, failures, and complexity notes.
3. **Reproducible inference.** The winner is frozen as relational operations, inferred temporal periods, an ordered feature schema, and a fitted estimator. Inference replays this artifact without rediscovering features.

Figure 1 summarizes the system. The relational database and point-in-time task contract enter a search over graph configurations. Each configuration changes the reachable graph and feature policy, producing a distinct feature-frame shape. A downstream learner scores that shape, KurverSC selects one program, and the chosen program is replayed for final fitting and prediction.




Figure 1: KurverSC system overview. GraphReduce materializes candidate relational representations; a downstream learner provides the validation objective; the selected graph program and schema are frozen for production fitting.

The rest of this report is organized as follows. Section 2 formulates relational signal compression. Section 3 describes the GraphReduce algorithm. Section 4 defines the feature families and annotations. Section 5 presents the KurverSC search and selection procedure. Section 6 describes final fitting and the fitted artifact. Section 7 covers systems behavior. Section 8 specifies the benchmark protocol and reserves the performance tables. Sections 9 and 10 discuss related work, limitations, and future directions.

5 Relational signal compression

5.1 Database, task, and prediction grain

Let a relational database be a directed schema graph

$$\mathcal{D} = (\mathcal{V}, \mathcal{E}),$$

where each node $v \in \mathcal{V}$ is a table and each edge $e \in \mathcal{E}$ declares a parent key, a child foreign key, and whether the child must be reduced before joining. A task specifies:

- a root or parent table r whose key defines the prediction grain;
- labeled entities and a target y ;
- training, validation, and test cutoffs or splits;
- date keys for tables whose rows are time-dependent; and
- a classification or regression metric.

For entity i and cutoff c_i , the eligible history is

$$\mathcal{H}_i(c_i) = \{x \in \mathcal{D} : c_i - H < t(x) \leq c_i\},$$

where $t(x)$ is a row timestamp and H is the configured compute horizon. The upper-bound convention is implemented explicitly at the KurveRSC root boundary; GraphReduce’s native feature windows remain strict internally where required. No row after the entity cutoff may contribute to its features.

5.2 A graph configuration is a representation hypothesis

KurveRSC represents a relational program by

$$q = (d, \mathcal{A}, a, b, x),$$

where:

- d is the maximum backward graph depth;
- $\mathcal{A}$ is the ordered, cumulative set of feature families;
- $a \in \{0, 1\}$ controls automatic annotation;
- b is the per-family source-column or predicate budget; and
- x contains optional text-annotation policy.

Executing q at cutoff c produces

$$F_q(c) = \text{GraphReduce}(\mathcal{D}, r, c, H; q) \in \mathbb{R}^{n_q(c) \times p_q(c)}.$$

The root key fixes the intended row grain. The feature width and lineage depend on the tables reachable within d , their types and values, the selected families, annotations, inferred lookbacks, and the budget b . In practice, entity eligibility can also change the observed row count. KurveRSC therefore records both n_q and p_q for every trial.

We call this *relational signal compression* because a variable-size collection of related histories is compiled into a bounded entity vector. Compression here does not mean a learned low-dimensional embedding. It means a deterministic, lineage-preserving reduction from a multi-table history to a tabular representation at the target grain.

5.3 The learner closes the objective loop

Let λ denote a downstream learner policy (backend and fixed hyperparameters). For candidate q ,

$$\theta_q = \text{Fit}_\lambda(F_q^{\text{train}}, y^{\text{train}}),$$

and

$$J(q; \lambda) = \text{Metric}(\theta_q, F_q^{\text{validation}}, y^{\text{validation}}).$$

The raw winner is




Figure 2: Candidate programs create different feature-frame shapes. Depth controls relational reach; feature families and annotations control the statistics preserved along each reached path. Validation selects a shape, after which the same relational program is materialized on the full production frame.

$$q_{\text{raw}} = \begin{cases} \arg \max_q J(q; \lambda), & \text{binary classification AUROC,} \\ \arg \min_q J(q; \lambda), & \text{regression MAE.} \end{cases}$$

This is a joint empirical objective over representation and learner behavior: the score is meaningful only after the learner has consumed the generated frame. It is not end-to-end differentiation through SQL. In the default implementation, λ is held fixed while q varies, which isolates relational representation quality. The API permits a different backend or caller-supplied model parameters; a broader AutoML layer could optimize (q, λ) jointly in future work.

6 GraphReduce execution algorithm

6.1 Graph representation and directional reduction

GraphReduce models tables as nodes and relationships as directed edges [1]. A node owns its source, primary key, optional date key, prefix, and table-local operations. An edge declares how the child key maps to the parent key and whether the child must be aggregated before the join.

The central invariant is grain preservation. For a one-to-many relationship, a raw child join would duplicate the parent. GraphReduce instead visits the deepest reachable children first, reduces each child by the relationship key, joins the result into its parent, and repeats toward the root. Deep descendants are thus compressed before their information propagates upward.

For a path $r \leftarrow v_1 \leftarrow \dots \leftarrow v_k$, let $R_{v_j \rightarrow v_{j-1}}$ denote the key-grouped reduction of node v_j . The propagated representation is recursively

$$Z_{v_{j-1}} = X_{v_{j-1}} \bowtie R_{v_j \rightarrow v_{j-1}}(Z_{v_j}),$$




Figure 3: GraphReduce performs a granularity-directed depth-first reduction. Dated rows are bounded by the entity cutoff and compute horizon. Children are annotated and reduced by their foreign key before the resulting statistics are joined toward the root grain.

and the final frame is Z_r . Prefixes and deterministic aliases retain column lineage across repeated aggregation.

6.2 Execution phases

At a high level, GraphReduce executes the following program:

Input: graph G , root r , cutoff c , compute horizon H ,
feature policy A , optional frozen plan P

1. Hydrate graph-level parameters onto nodes.
2. Load each table and assign collision-safe prefixes.
3. Annotate node-local values and the dated root's age at cutoff c .
4. Apply node filters and normalization operations.
5. Propagate entity-specific cutoff context through reachable relationships.
6. Traverse reducible edges deepest-child first:
 - a. restrict dated child rows to $(c - H, c]$;
 - b. infer or replay annotation and feature operations;
 - c. group by the child foreign key and aggregate;
 - d. join the reduced child into its parent.
7. If training labels are graph-generated, execute the future-label window separately through GraphReduce's label operations.
8. Return the root-grain frame and, during planning, the executed SQL plan.

The same abstraction supports custom node operations. KurveRSC deliberately uses GraphReduce’s native `prep_for_features()`, `prep_for_labels()`, `do_reduce()`, and `do_labels()` lifecycle rather than inventing a second label-generation system. External benchmarks such as RelBench supply labels and official cutoffs; native GraphReduce tasks may instead generate labels from a connected label node.

6.3 Point-in-time correctness

Temporal correctness is enforced at multiple boundaries:

1. **Root filtering.** When the parent has a date key, KurveRSC inserts `date <= cutoff` into both the root load and GraphReduce filtering stages. This prevents post-cutoff parent rows from reaching schema inference, annotation, or joins.
2. **Relationship windows.** A dated child is filtered to the configured compute horizon before reduction. If entity-specific cutoff rows are used, the cutoff is propagated down the graph so every related event is compared against the correct entity timestamp.
3. **Label separation.** Feature history ends at the cutoff. Graph-generated labels use GraphReduce’s label period after the cutoff, while external labels follow the benchmark’s official task construction.
4. **Plan replay.** Cutoff-dependent literals in a frozen plan are rebound to the new cutoff while the learned operations and schema remain unchanged.
5. **Test isolation.** Test targets do not participate in graph configuration, annotation discovery, inferred periods, schema selection, or model fitting.

RelBench defines official database truncation, temporal splits, task labels, and evaluation contracts for RelBenchV1 [4]. The KurveRSC RelBench adapter consumes those production objects instead of reconstructing task semantics from local assumptions.

6.4 Automatic event-frequency inference

A fixed lookback grid is not equally appropriate for daily clicks, monthly orders, and sparse study events. With `infer_ts_periods=True`, GraphReduce infers lookback periods per dated node and relationship. For each relationship key it computes first timestamp, last timestamp, and event count, then estimates an average cadence from duration and count.

Let D be mean observed active duration in days, N mean event count, and H the actual graph compute horizon. GraphReduce builds candidate windows from ratios of D and H , retains cadence and horizon anchors, rounds long windows to stable calendar-like intervals, and selects between five and ten distinct periods depending on duration and event density. Every inferred period is bounded by H ; the final grid always spreads short-, medium-, and long-range windows over the actual compute horizon.

The selected periods become part of the frozen execution plan. Validation, historical replay, and test inference restore them rather than re-estimating them from later data.

6.5 Frozen execution plans

Automatic feature generation can depend on observed types, cardinalities, frequent categories, or available annotations. Rediscovering these choices at each cutoff could silently change feature meaning. GraphReduce therefore records the executed operations, date-filter operations, reduction keys, inferred periods, and original planning cutoff. KurveRSC augments the plan with the ordered model feature schema and categorical/datetime roles.

For a later cutoff c' , plan replay:

- preserves the exact selected operations and aliases;
- rebinds cutoff and lookback literals from c to c' ;
- restores inferred `ts_periods`;
- prohibits feature rediscovery; and
- aligns the resulting frame to the training-derived schema.

A deterministic fingerprint identifies the production plan stored in the fitted artifact.

7 Relational feature construction

7.1 Type-aware base reductions

GraphReduce infers semantic and physical types from a bounded sample, excludes keys from inappropriate aggregation, and emits SQL aggregations compatible with the selected backend. The conservative base behavior includes:

- row or identifier counts;

- numerical `sum`, `avg`, `min`, and `max` where type-safe;
- Boolean sums;
- bounded categorical counts and frequent-value summaries;
- time since the last event;
- event counts inside configured windows and adjacent-window change ratios; and
- bounded frequent predicates with windowed counts, lag-bucket counts, and time since the last matching event.

GraphReduce avoids numeric SQL functions on physically non-numeric values, even when a sample-level semantic inference is ambiguous. This type guard is important for production relational data, where an all-null sample can otherwise be exposed by a SQL engine as a string-typed placeholder.

7.2 Automatic annotations

Annotations convert row-level values into aggregation-friendly signals before the child is reduced. When `auto_annotate_features=True`, GraphReduce can emit:

- indicators for bounded-cardinality categorical values;
- top-*k* indicators and an `other` bucket for higher-cardinality categories;
- category-gated numerical values, such as amount when status equals a frequent category;
- caller-provided predicate or value expressions; and
- optional text-shape features: character length, word count, number presence, URL presence, and empty-text indicators.

The dated graph parent receives an age-in-days feature relative to the cutoff, independent of whether generic automatic annotations are enabled. Annotation budgets prevent one wide source table from dominating the graph.

7.3 Feature families

GraphReduce implements seven feature families. KurveRSC searches independent subsets of the four general-purpose optional families rather than requiring a weak earlier family to remain in every later candidate. A forward beam controls the resulting combinatorial growth.

Family	Signal preserved	Principal operations
<code>base</code>	Stable type-aware summaries and first-order recency	Counts; numeric reductions; categorical summaries; recency; rolling counts and changes; bounded predicate histories
<code>temporal</code>	Magnitude and distribution of numerical behavior over time	Windowed sum, average, minimum, and maximum for selected numeric or annotated columns
<code>sequence</code>	Trajectory, concentration, and burstiness	Activity rate per day; recent-window share of lifetime activity; adjacent-window burst ratios; active span; activities per active day
<code>conditional</code>	Behavior under categorical or semantic conditions	Windowed conditional count, share, presence, and adjacent-window change
<code>episode</code>	Exposure and distinct-event opportunity	Lifetime and windowed row counts; distinct primary-key counts where supported
<code>semantic</code>	Caller-declared domain meaning	Predicate and value expressions compiled to annotated columns and consumed by ordinary reductions
<code>context</code>	Peer-relative row context before reduction	Peer-group size and bounded numeric deviations over caller-declared context keys

`semantic` and `context` remain available in GraphReduce but are omitted from the default KurveRSC frontier. Semantic expressions require domain declarations, and context features require meaningful context keys; neither can be assumed for a fully generalized task. Their signals can still be supplied through explicit `GraphConfig` candidates.

7.4 Feature budgets and controlled width

KurveRSC fits a deterministic source-column ordering on a small sample from each node. Structural join and cutoff columns are retained first; remaining columns are ranked by observed coverage, entropy, spread, and cardinality, with effectively unique identifiers down-weighted. This ordering replaces physical schema order as the input to the capped feature funnel and is returned as `feature_audit`.

Three independent width controls are frozen in each graph program:

1. `feature_family_max_columns=b` limits selected source columns per family and node;
2. `feature_family_max_features_per_column=r` limits the derived SQL features emitted by each selected source column within a family; and
3. `feature_propagation_max_functions_per_column=p` limits the functions emitted when an already-derived column crosses another reduction.

For $p = 1$, propagation retains the canonical associative or additive continuation: max-to-max, min-to-min, sum-to-sum, count-to-sum, and average/share-to-average. Thus a useful lineage such as `max(max(max(value)))` can cross three hops without branching into every compatible aggregation at every hop.

KurveRSC currently fixes several lower-level GraphReduce policies across graph trials so the main meta-parameter comparison remains interpretable:

GraphReduce policy	KurveRSC reference value
Initial lookback periods	7, 30, and 90 days
Categorical cardinality threshold	20
Frequent-category top- k	5
Maximum auto-annotated categorical columns	10
Maximum gated numeric columns	4
Categories per gated numeric column	3
Initial / promoted source-column budgets	4 / 8
Derived features per selected source column	32
Aggregate functions per propagated column	1
Text features	Disabled by default

When event-frequency inference is enabled, its learned periods replace the initial lookback grid for the affected node or relationship.

8 KurveRSC meta-optimization

8.1 Search dimensions

The default KurveRSC frontier independently varies four primary meta-parameters:

1. **Depth** $d \in \{1, 2, 3\}$, controlling how many backward relational hops may contribute to the root.
2. **Feature-family subset** $\mathcal{A}$, controlling which kinds of relational signal survive compression.
3. **Automatic annotation** $a \in \{\text{on}, \text{off}\}$, testing whether inferred categorical and gated-numeric predicates help the task.
4. **Source-column budget** b , fixed at four by default and optionally refined to a caller-requested wider tier for the strongest shapes.

The optional-family set is:

`base + any subset of {temporal, sequence, conditional, episode}`

Depth 3 is evaluated only when the family subset is contained in `base`, `temporal`, and `sequence`. Candidates

containing `conditional` or `episode` are limited to depths 1 and 2. With both annotation settings, one source-column budget has 72 potential graph shapes:

$$(4 \text{ compact subsets})(3 \text{ depths})(2 \text{ states}) + (12 \text{ wider subsets})(2 \text{ depths})(2 \text{ states}) = 72.$$

The narrow-budget forward beam materializes at most 24 of those shapes. Both annotation states are screened at base depth 1, only the stronger state is promoted to depth 2, and depth 3 requires a meaningful depth-2 gain. This admits at most four base variants; beam width two admits at most 8, 6, 4, and 2 descendants at the successive optional-family levels. The other shapes remain in the audit as pruned candidates. An optional second 72-shape lattice can represent an eight-column budget, but only the raw narrow-budget winner and its complexity-aware recommendation are promoted; the next score-ranked shape fills the second slot when those are identical. The default funnel therefore contains 72 auditable potential candidates and normally materializes no more than 24.

The user may change `max_depth`, family stages, annotation options, and feature budgets, or supply an explicit list of `GraphConfig` objects. Optional text features apply consistently to those candidates. Graph search is therefore bounded by default but not hard-coded to one benchmark.

8.2 Screening protocol

Each configuration receives a fresh `GraphReduce` instance because graph execution creates intermediate state. A screening trial performs the following steps:

```
rank source columns and freeze the feature-funnel audit
for q in forward_beam(candidate_graph_configs):
    build a fresh depth-bounded graph for q
    materialize the latest eligible training cutoff
    learn and freeze q's GraphReduce operations and feature schema
    fit the downstream learner on that training frame
    stream validation cutoff frames through the frozen plan
    aggregate predictions and compute AUROC or MAE
    record score, standard error, rows, columns, timings, plan, and failures
promote only the strongest complete shapes to the wider column budget
```

The report’s default reproducibility profile sets `search_full_data=True`: every beam-admitted graph configuration is evaluated against the complete source tables at the latest eligible training cutoff. The strongest three configurations are then reranked on three sequential complete relational cutoff folds. This is the configuration used for the primary result tables unless a row is explicitly identified as an ablation or retry.

For lower-cost exploratory runs, `search_full_data=False` enables the public API’s three-fidelity funnel. Beam-admitted candidates are screened with `screening_rows=10000`; `confirmation_top_k=8` structurally diverse candidates are rebuilt with `sample_rows`; and the strongest `rerank_top_k` candidates are reranked on complete folds. Every mode searches exactly one training cutoff per node at a time and executes candidates sequentially.

Reference-stage choice	Parameterized setting
Complete latest-cutoff graph search	<code>search_full_data=True</code>
Sequential graph candidates	intrinsic per-task execution policy
Full-data finalists	<code>rerank_top_k=3</code>
Sequential stability folds	<code>rerank_cutoff_frames=3</code>
Final production cutoffs	<code>full_training_frames=1</code>
Inferred event periods	<code>infer_ts_periods=False</code>
Automatic text annotations	<code>auto_text_features=False</code>

Validation cutoffs are materialized one at a time. Once the learner has produced predictions, the frame is released. This preserves the meaning of the metric while bounding resident feature-frame memory.

8.3 Downstream learner policy

CatBoost is the current default because it provides a strong conventional model for mixed numerical and categorical frames and handles categorical features natively [5]. The reference policy uses:




Figure 4: Default RelArena reproducibility profile. Every admitted graph configuration is evaluated on the complete latest-cutoff source data. The top three candidates are reranked over three complete cutoff folds processed sequentially; the stability-adjusted winner is frozen and fitted on one complete latest production cutoff before test replay.

Parameter	Value
Trees	300
Depth	6
Learning rate	0.05
L2 leaf regularization	10
Native threads per task	8
Classification loss / selection metric	Log loss / AUROC
Regression loss / selection metric	MAE / MAE

Caller-supplied `model_params` override this policy. Those parameters are held constant across relational candidates in a run. This is important: KurveRSC attributes score changes to the relational configuration unless the caller explicitly constructs a larger joint search.

The estimator boundary is pluggable. An experimental `tabpfn_v3` backend exists for single-frame fitting, but streaming multi-frame production fitting currently requires CatBoost. TabPFN and its successors motivate treating relational flattening as a separate program from tabular inference [7–9]; KurveRSC exposes exactly that separation.

8.4 Complexity-aware recommendation

The raw best validation score can favor a much wider representation for a negligible gain. KurveRSC therefore retains two results:

- `best_config`: the raw metric winner; and
- `recommended_config`: the simplest configuration within a declared tolerance of that winner.

KurveRSC estimates validation uncertainty for every candidate: an analytic standard error for AUROC and the standard error of absolute residuals for MAE. For classification, candidate q is eligible when

$$J(q) \geq J(q_{\text{raw}}) - \max(\epsilon_{\text{AUC}}, m\sqrt{\text{se}(q_{\text{raw}})^2 + \text{se}(q)^2}),$$

with default $\epsilon_{\text{AUC}} = 0.002$ and uncertainty multiplier $m = 1$. For regression, it is eligible when

$$J(q) - J(q_{\text{raw}}) \leq \max(|J(q_{\text{raw}})|\epsilon_{\text{MAE}}, m\sqrt{\text{se}(q_{\text{raw}})^2 + \text{se}(q)^2}),$$

with default $\epsilon_{\text{MAE}} = 0.005$.

Among eligible candidates, KurverSC minimizes feature count, then nominal graph complexity, then materialization time. The nominal complexity is

$$C(q) = d|\mathcal{A}|(1 + 0.25a).$$

A separate note flags cases where feature width grows by at least the configured factor (2x by default) without a meaningful metric gain. This does not erase the raw winner; it makes the cost of choosing it explicit and allows the recommended configuration to resist overfitting from representational expansion. Setting `complexity_uncertainty_multiplier=0` recovers selection based only on the fixed practical tolerances.

8.5 Full-frame and optional walk-forward stability reranking

One-cutoff screening may favor a configuration whose signal is unstable over time. With `rerank_top_k > 0`, KurverSC evaluates a bounded top- k finalist set on full-data cutoff folds. The finalist budget contains the strongest raw medium-fidelity candidates. Setting `rerank_cutoff_frames=1` evaluates the finalists on only the official full-data train-to-validation frame.

For candidate q , fold scores $s_1, \dots, s_K$ are summarized as

$$\bar{s}_q = \frac{1}{K} \sum_{k=1}^K s_k, \quad \tilde{J}(q) = \begin{cases} \bar{s}_q - \gamma \text{sd}(s_q), & \text{AUROC,} \\ -\bar{s}_q - \gamma \text{sd}(s_q), & \text{MAE,} \end{cases}$$

where γ is `rerank_stability_penalty` (0.25 by default when reranking is enabled). Each fold learns only from its training cutoff and scores the next cutoff or official validation cutoffs. One frame is resident at a time. The raw maximizer of $\tilde{J}$ is the selected production configuration; complexity annotations remain available for audit and earlier-fidelity promotion but do not override the full-frame result.

The default evaluates three finalists on three sequential full-data folds (`rerank_top_k=3, rerank_cutoff_frames=3`).

8.6 Resource-aware failure handling

Relational feature spaces can exceed memory or temporary-disk budgets. KurverSC records candidate failures instead of losing the experiment. When a candidate fails with a recognized out-of-memory condition, deeper candidates with the same family prefix and annotation state are resource supersets; they may be skipped with an explicit audit record. This prevents a known-infeasible branch from repeatedly exhausting the host.

9 Final fitting and prediction

9.1 Selecting a shape and freezing its meaning

The recommended candidate after screening, or after optional reranking, is the production graph configuration q^* . The fitted relational component is not only the tuple $(d, \mathcal{A}, a, b)$. It includes:

- the normalized table and relationship specification;
- GraphReduce operations selected from observed training data;
- inferred temporal periods per node and relationship;
- cutoff-rebindable temporal predicates;
- ordered feature names and semantic roles; and
- a plan fingerprint.

This distinction matters because automatic operations can depend on training cardinalities and frequent values. Saving only the top-level configuration would not guarantee that a later graph reconstructs the same columns.

9.2 Multiple production frames

`full_training_frames=K` controls how many eligible training cutoffs are used for the final artifact. If $K = 1$, KurveRSC uses the latest eligible training cutoff. For $K > 1$, it materializes one cutoff frame, fits one independent CatBoost model, releases the frame, and proceeds to the next cutoff. Prediction averages the models:

$$\hat{f}_{\text{train}}(x) = \frac{1}{K} \sum_{k=1}^K f_k(x).$$

This frame ensemble avoids retaining K large matrices in RAM and avoids the ordering sensitivity observed when appending trees incrementally across temporal frames. Each member receives the same 300-tree reference budget unless the user overrides it.

The default reproducibility profile sets `full_training_frames=1`. Experiments with three or more cutoff frames have improved held-out performance on some tasks by exposing the learner to additional temporal regimes. They are not the default because they multiply full graph replays, learner fits, elapsed time, spill exposure, and the fitted ensemble size. Sequential execution still keeps peak feature-frame residency at one frame, but the one-cutoff default is the most portable configuration for automated RelArena reproduction. Callers can set any positive K ; selected cutoffs are deterministically spread across the available training period, while `None` uses all eligible cutoffs.

The validation ensemble consists only of training-frame models. After validation is scored, independent models fitted on validation frames may be added to the final train-plus-validation ensemble. Test prediction uses that final ensemble; validation labels never influence the model that produced the reported validation score.

9.3 One-frame resident lifecycle

The production loop is deliberately streaming:

```
select q* and frozen plan P*

training_models = []
for cutoff in selected_training_cutoffs:
    frame = replay GraphReduce(P*, cutoff)
    model = fit learner(frame)
    training_models.append(model)
    release frame

score validation cutoffs with ensemble(training_models)

for cutoff in validation_cutoffs:
    frame = replay GraphReduce(P*, cutoff)
    validation_models.append(fit learner(frame))
    release frame

final_model = ensemble(training_models + validation_models)
replay P* with train=False for test or future cutoffs
```

At no point does the algorithm require several cutoff feature frames to coexist. The main residual memory costs are the active GraphReduce/DuckDB materialization, DuckDB buffers and temporary data, one pandas frame crossing the learner boundary, and the comparatively small fitted model ensemble.

9.4 Fitted artifact

The returned `FitResult` exposes the complete trial audit and a `FittedKurveRSC` artifact. The artifact includes the selected `GraphConfig`, frozen plan, feature schema, downstream estimator or frame ensemble, task metadata, target classes, validation/test scores when available, and prediction helpers. A preselected configuration and plan can be supplied to bypass search and reproduce a known relational representation.

10 Systems implementation

10.1 SQL execution with DuckDB

KurveRSC’s reference path instantiates GraphReduce with DuckDB. DuckDB is an in-process analytical database designed for vectorized analytical execution [6]. Source tables are registered as views, GraphReduce emits SQL for annotation, filtering, reduction, and joining, and the final root frame crosses into pandas only at the downstream learner boundary.

The implementation constrains DuckDB memory and temporary storage per task, and limits CatBoost’s native threads. These guardrails are particularly important when RelArena launches many tasks concurrently: task-level parallelism multiplies both query buffers and native model workers.

The limits are not tied to one machine. `duckdb_memory_limit` and `duckdb_max_temp_directory_size` parameterize the per-task DuckDB memory and spill ceilings. A benchmark retry may raise them for an unusually wide relational join, but must report the override. They do not change the selected search space or permit several feature frames from the same task to coexist.

10.2 Search and task parallelism

Graph configurations within one task execute sequentially. This avoids concurrent construction of multiple large frames for the same database. RelArena may instead parallelize independent tasks through `--parallel-tasks N`. Frame construction inside a task remains serial and bounded by its DuckDB connection.

The two controls solve different problems:

- `parallel_tasks` governs throughput across benchmark tasks;
- `full_training_frames` governs temporal coverage inside one fitted artifact.

Within a task, `rerank_cutoff_frames` controls how many temporal confirmation folds are replayed sequentially, while `rerank_top_k` controls how many candidate shapes receive that confirmation. `search_full_data` switches graph-configuration selection between complete and sampled source tables. These controls alter total work without changing the one-resident-frame invariant.

Neither changes graph depth, feature-family width, or the number of graph candidates.

10.3 Reproducibility and observability

Structured logs report search start, each candidate configuration, candidate completion or failure, rows, features, materialization time, model time, selection, reranking, production-plan fingerprint, refit, test prediction, and per-task peak resident memory. The trial table is part of the result object.

For a reproducible report, a run should disclose:

- KurveRSC, GraphReduce, RelArena, RelBench, DuckDB, and learner versions;
- the exact candidate space and feature budgets;
- sampling versus full-data search;
- inferred-period and text-feature settings;
- number of production cutoff frames;
- model parameters and random seed;
- per-task resource ceilings and task concurrency; and
- hardware, wall-clock time, peak RSS, failures, and retries.

11 Experimental protocol and results

11.1 Research questions

The controlled evaluation is designed to answer:

1. Does learner-guided graph selection outperform a fixed shallow relational representation?
2. Which feature-family stages and depths are selected across tasks?
3. How often does complexity-aware recommendation preserve performance while reducing width or runtime?
4. Does inferred event frequency improve over a fixed temporal grid?

5. What is the value of additional full-training cutoff frames?
6. What are the runtime and peak-memory costs relative to relational baselines?
7. How sensitive are selected configurations to downstream learner choice?

11.2 Benchmark and evaluation discipline

The primary benchmark is RelBenchV1 entity classification and regression [4], executed through RelArena. RelArena provides a common model interface and cross-task comparison context [10]. The evaluation must use the production RelBench dataset implementation, official task objects, timestamps, train/ validation/test splits, and evaluators.

Binary classification is reported with test AUROC. Regression is reported with test MAE, where lower is better. Aggregate reporting should include mean rank, pairwise win rates, and the exact RelArena leaderboard statistic used by the corresponding RelArena release. Validation results are used only to select the graph configuration; all headline results use held-out test labels.

11.3 Main classification results

The following tables are the September 1, 2026 snapshot of the default reproducibility profile. KurveRSC uses full-data latest-cutoff graph search, three finalists, three sequential reranking folds, one production training cutoff, CatBoost, fixed GraphReduce temporal periods, and no automatic text annotations. Every other value is the validation-selected, seed-zero test score from RelArena’s reproduced release artifact. These are raw held-out scores, not normalized leaderboard values. Bold marks the best result across all included models and systems, and the final column names that overall winner.

The compact PDF headings use **TPFN API/Local** for TabPFN-Rel API/Local, **G-SAGE** for GraphSAGE, **RDB** for RDBLearn, **LGBM** for LightGBM, and **Const/E/G** for the per-entity/global constant predictors.

Dataset / task	Kurve	RT-P	TPFN API	TPFN Local	G-SAGE	RelGT	RDB	RelGNN	LGBM	Const/E	Overall Const/G winner
Amazon / item-churn	0.828114	0.832656	0.827996	0.827857	0.830527	0.823835	0.819538	0.785579	0.662211	0.728875	0.500000 RT-PluRel
Amazon / user-churn	0.709790	0.713460	0.708649	0.702403	0.704596	0.701924	0.684375	0.694281	0.517056	0.634205	0.500000 RT-PluRel
Avito / user-clicks	0.654236	0.583376	0.675191	0.614522	0.608674	0.644353	0.678769	0.667571	0.564163	0.504143	0.500000 RDBLearn
Avito / user-visits	0.674295	0.670887	0.668026	0.668811	0.665758	0.662142	0.659624	0.648731	0.529282	0.602703	0.500000 KurveRSC
Event / user-ignore	0.830661	0.847577	0.878659	0.701376	0.758728	0.781507	0.664351	0.805393	0.777181	0.839930	0.500000 TPFN API
Event / user-repeat	0.780414	0.791377	0.759291	0.769251	0.784626	0.734358	0.744084	0.754612	0.748295	0.751805	0.500000 RT-PluRel
F1 / driver-dnf	0.753628	0.731460	0.732172	0.714468	0.717235	0.711667	0.714551	0.726106	0.730298	0.699258	0.500000 KurveRSC
F1 / driver-top3	0.673136	0.758858	0.771426	0.792916	0.725975	0.810841	0.780081	0.758864	0.738889	0.556530	0.500000 RelGT
H&M / user-churn	0.696866	0.704356	0.705215	0.705690	0.698525	0.689531	0.698352	0.682025	0.590081	0.647972	0.500000 TPFN Local
Stack / user-badge	0.877852	0.891612	0.880386	0.863470	0.888748	0.574286	0.771147	0.620584	0.537995	0.788956	0.500000 RT-PluRel
Stack / user-engagement	0.903450	0.896775	0.905994	0.905834	0.905609	0.906731	0.858670	0.905054	0.811836	0.826717	0.500000 RelGT
Trial / study-outcome	0.704273	0.723487	0.764702	0.730607	0.686232	0.668495	0.721205	0.657435	0.715018	0.500000	0.500000 TPFN API

The isolated `rel-amazon/user-churn` retry completed with validation AUROC 0.704286 and test AUROC 0.709790. It selected a depth-two ("`base`", "`sequence`") graph with automatic annotation disabled and 135 materialized features. Peak resident memory was 79.4 GiB.

11.4 Main regression results

Dataset / task	Kurve	RT-P	TPFN API	TPFN Local	G-SAGE	RelGT	RDB	RelGNN	LGBM	Const/E	Const/G	Overall winner
Amazon / item-ltv	42.205991	43.028012	46.768181	47.768328	49.245337	48.673386	48.997527	52.458084	55.750750	65.351419	64.233613	KurveRSC
Amazon / user-ltv	14.132556	13.943011	14.358212	14.400940	14.415321	14.352129	14.577540	14.575225	16.784682	17.423469	16.782979	RT-PluRel
Avito / ad-ctr	0.033654	0.034848	0.031080	0.031379	0.038966	0.036499	0.034103	0.042646	0.041250	0.041201	0.043067	TPFN API
Event / user-attendance	0.258185	0.240949	0.243944	0.239383	0.245018	0.261493	0.242215	0.243858	0.262798	0.269152	0.263534	TPFN Local
F1 / driver-position	3.913762	3.817699	3.769181	3.761699	4.011112	4.765529	3.888649	4.265887	4.105829	4.103509	4.399101	TPFN Local
H&M / item-sales	0.032107	0.040258	0.060507	0.061362	0.055153	0.053168	0.067132	0.056493	0.075295	0.078033	0.076087	KurveRSC
Stack / post-votes	0.063347	0.063463	0.067882	0.067957	0.064898	0.067904	0.067719	0.067904	0.066099	0.069374	0.067904	KurveRSC
Trial / site-success	0.402717	0.410357	0.412624	0.385751	0.324851	0.370110	0.485833	0.340493	0.437506	0.441148	0.462222	GraphSAGE
Trial / study-adverse	42.024547	32.652791	39.753674	42.591708	44.315261	44.125887	44.026705	46.270064	44.573495	57.533247	57.533247	RT-PluRel

Across all 21 direct comparisons, KurveRSC wins 13 and TabPFN-Rel Local wins 8: 8–4 on classification and 5–4 on regression.

11.5 Aggregate Elo leaderboard

RelArena requires a dense method-by-task matrix. The completed KurveRSC result set supplies all 21 tasks, so the unmodified `compute_leaderboard` procedure can fit the complete matrix without dropping a task. It uses `bencheval`’s Bradley–Terry fit, anchors the global constant predictor at 1000, and bootstraps uncertainty with seed zero. The table includes every reproduced release method, including the entity-only LightGBM baseline that RelArena omits from its compact README display. It is ordered by Elo, not by normalized loss.

Elo rank	Method	Kind	Elo	Bootstrap –/+	Mean rank	Win rate	Rescaled loss
1	RT-PluRel	system	1858.7	–102.9/ +92.3	2.952	80.48%	0.107752
2	TabPFN-Rel API	model	1824.0	–67.5/ +88.7	3.286	77.14%	0.148825
3	KurveRSC	system	1786.7	–88.1/ +88.8	3.667	73.33%	0.142977
4	TabPFN-Rel Local	model	1725.5	–87.1/ +88.4	4.333	66.67%	0.190129
5	GraphSAGE	model	1663.0	–75.3/ +83.3	5.048	59.52%	0.207151
6	RelGT	model	1578.7	–106.5/ +123.8	6.024	49.76%	0.321981
7	RDBLearn	model	1564.1	–101.6/ +90.9	6.190	48.10%	0.274366
8	RelGNN-ES	model	1528.0	–78.0/ +89.8	6.595	44.05%	0.306010
9	LightGBM (entity-only)	model	1359.8	–126.6/ +83.0	8.286	27.14%	0.546964
10	Constant (per-entity)	model	1259.5	–147.3/ +124.3	9.095	19.05%	0.618485
11	Constant (global)	model	1000.0	–139.6/ +74.5	10.524	4.76%	0.941953

11.5.1 Ranking by rescaled loss

RelArena’s default aggregate ranking uses mean per-task min-max rescaled error, where lower is better. The following table presents that ordering separately from Elo and retains each method’s Elo rank for comparison.

Loss rank	Method	Kind	Rescaled loss	Elo rank	Elo
1	RT-PluRel	system	0.107752	1	1858.7
2	KurveRSC	system	0.142977	3	1786.7
3	TabPFN-Rel API	model	0.148825	2	1824.0
4	TabPFN-Rel Local	model	0.190129	4	1725.5
5	GraphSAGE	model	0.207151	5	1663.0
6	RDBLearn	model	0.274366	7	1564.1
7	RelGNN-ES	model	0.306010	8	1528.0
8	RelGT	model	0.321981	6	1578.7
9	LightGBM (entity-only)	model	0.546964	9	1359.8
10	Constant (per-entity)	model	0.618485	10	1259.5
11	Constant (global)	model	0.941953	11	1000.0

On Elo, KurveRSC is third overall and second among the two system submissions. On RelArena’s default `loss_rescaled` ordering it is second overall (0.142977), behind RT-PluRel (0.107752) and immediately ahead of TabPFN-Rel API (0.148825).

11.6 Configuration-selection summary

Quantity	Classification	Regression	Overall
Tasks	12	9	21
Mean successful candidates	—	—	—
Mean selected depth	—	—	—
Annotation enabled selections	—	—	—
Mean selected feature count	—	—	—
Complexity recommendation differs from raw winner	—	—	—
Inferred-period selections	—	—	—

11.7 Ablations

Variant	Classification aggregate	Regression aggregate	Mean feature count	Mean runtime	Peak RSS
Full	—	—	—	—	—
KurveRSC	—	—	—	—	—
Fixed depth	—	—	—	—	—
1 + base	—	—	—	—	—
No	—	—	—	—	—
automatic	—	—	—	—	—
annotations	—	—	—	—	—
Fixed	—	—	—	—	—
7/30/90-day	—	—	—	—	—
periods	—	—	—	—	—
No	—	—	—	—	—
complexity	—	—	—	—	—
recommenda-	—	—	—	—	—
tion	—	—	—	—	—
One	—	—	—	—	—
production	—	—	—	—	—
frame	—	—	—	—	—
Three	—	—	—	—	—
production	—	—	—	—	—
frames	—	—	—	—	—
Walk-	—	—	—	—	—
forward	—	—	—	—	—
finalist	—	—	—	—	—
reranking	—	—	—	—	—

11.8 Runtime and memory

Runtime should separate graph search, optional reranking, final graph replay, learner fitting, and prediction. Peak RSS must be reported per task and for the parallel sweep host because aggregate parallel memory is not the same as a single-model requirement.

Model	Data preparation	Model fit	Prediction	End-to-end	Peak RSS	Hardware
KurveRSC	—	—	—	—	—	—

Model	Data preparation	Model fit	Prediction	End-to-end	Peak RSS	Hardware
TabPFN-Rel Local	—	—	—	—	—	—
RT-PluRel	—	—	—	—	—	—

No external baseline runtime should be inferred from predictive-score tables. If a method does not publish end-to-end relational preparation time or peak memory, the table should say “not reported” rather than estimate it.

12 Related work

12.1 Automated relational feature engineering

Deep Feature Synthesis recursively composes transformation and aggregation primitives over entity relationships [2]. OneBM represents relational paths and feature candidates in a scalable automated system [3]. GraphReduce shares the goal of automatic relational feature construction but emphasizes an explicit compute graph, depth-first grain-preserving reduction, customizable node operations, SQL execution, and reusable frozen plans [1]. KurveRSC adds an outer validation loop that chooses among GraphReduce programs using a downstream learner.

12.2 Relational deep learning and foundation models

RelBench standardizes relational datasets, temporal tasks, and evaluation [4]. GraphSAGE learns inductive node representations through neighborhood aggregation [11]. RelGNN develops composite message passing for relational databases [12], and RelGT applies relational graph transformers [13]. These systems learn through graph neural architectures rather than compiling a database into a conventional tabular frame.

RDBLearn shows that tabular foundation models can make predictions over relational databases after automatic flattening [9]. TabPFN-3 reports relational results through TabPFN-Rel and highlights the sensitivity of relational methods to data regime and evaluation protocol [8]. KurveRSC is complementary: it treats the flattening program itself as the meta-parameter space and permits CatBoost, TabPFN, or another learner to evaluate that space.

12.3 Tabular learning

CatBoost reduces target leakage in categorical boosting and remains a strong mixed-type tabular baseline [5]. TabPFN frames supervised tabular prediction as in-context inference under a learned synthetic prior [7], and TabPFN-3 scales that paradigm to substantially larger row and feature regimes [8]. KurveRSC does not replace either family. It produces a point-in-time relational representation and uses the selected learner as the empirical judge of that representation.

13 Limitations and future work

KurveRSC’s present limitations define a concrete research agenda:

- **Finite search.** The candidate lattice and beam are deterministic. They do not continuously optimize depth, windows, budgets, and learner parameters.
- **Learner-conditioned selection.** A graph that is optimal for CatBoost may not be optimal for a tabular foundation model. This is expected, not a violation of the objective; the representation is selected for the declared learner policy.
- **One-cutoff broad search.** The default full-data graph search sees the latest training cutoff; the top three candidates are then checked over three walk-forward folds. Earlier temporal regimes can still contain signal that is not represented in the broad candidate ordering.
- **One-cutoff production default.** Fitting on additional evenly spaced training cutoffs can improve some tasks, but increases graph replay, learner fitting, and artifact cost. It is exposed through `full_training_frames` rather than enabled in the bounded-memory reproducibility profile.
- **Binary classification and regression.** The reference metric path currently targets RelBenchV1 entity binary classification and regression. Multiclass, ranking, survival, and forecasting objectives require explicit adapters.

- **Heuristic feature budgets.** Utility-ranked source, per-column expansion, and propagation budgets constrain multiplicative growth but do not directly optimize memory, latency, or energy.
- **Text is shallow by default.** Optional text annotations capture shape and surface indicators, not semantic embeddings. A text-aware learner or embedding stage could recover additional signal.
- **Schema quality matters.** Missing or incorrect keys and date columns cannot be repaired by configuration search. Composite and many-to-many relationships require careful task adaptation.
- **Execution cost.** Full-frame evaluation of roughly 28 beam-admitted graph programs is substantially more expensive than fitting one fixed representation. Adaptive pruning, successive halving, and learned candidate priors are natural extensions.

Future work includes cost-aware multi-objective search, adaptive family staging, learner-specific graph priors, richer semantic annotations, native multiclass metrics, distributed execution, plan caching across related tasks, and systematic comparison of CatBoost and tabular foundation models on the same frozen relational frames.

14 Conclusion

KurveRSC turns relational representation from a fixed preprocessing decision into a validation-tested component of model fitting. GraphReduce provides the grain-preserving execution substrate: it filters history at entity cutoffs, annotates rows, recursively reduces child tables, propagates features toward a root entity, and records a replayable SQL plan. KurveRSC proposes alternative depths, feature-family stages, annotation policies, and feature budgets; fits a downstream learner on each resulting frame; and keeps the representation whose utility survives held-out evaluation and a complexity check.

The selected “shape” is more than a width. It is a relational program with table reach, temporal semantics, feature lineage, inferred windows, and an ordered schema. Freezing that program makes final fitting and future prediction reproducible. Streaming one frame at a time makes multiple temporal training cutoffs possible without multiplying resident feature-frame memory.

CatBoost is the current reference learner, not a permanent architectural constraint. The enduring abstraction is the coupled loop:

propose relational program $\rightarrow$ materialize frame $\rightarrow$ fit learner $\rightarrow$ validate $\rightarrow$ freeze and replay winner.

This loop is KurveRSC’s core contribution: an optimization over the relational signal-compression space, grounded in the predictive behavior of the learner that will ultimately use it.

15 References

1. W. Madrigal. **GraphReduce: scalable feature engineering with graphs of tables, relationships, and compute operations.** Open-source software and technical report, Kurve, Inc. [GitHub](#), [documentation](#), and [PyPI](#).
2. J. M. Kanter and K. Veeramachaneni. **Deep Feature Synthesis: Towards Automating Data Science Endeavors.** IEEE International Conference on Data Science and Advanced Analytics, 2015. [doi:10.1109/DSAA.2015.7344858](#).
3. H. T. Lam, J.-M. Thiebaut, M. Sinn, B. Chen, T. Mai, and O. Alkan. **One Button Machine for Automating Feature Engineering in Relational Databases.** 2017. [arXiv:1706.00327](#).
4. J. Robinson, R. Ranjan, W. Hu, K. Huang, J. Han, A. Dobles, M. Fey, J. E. Lenssen, Y. Yuan, Z. Zhang, X. He, and J. Leskovec. **RelBench: A Benchmark for Deep Learning on Relational Databases.** NeurIPS Datasets and Benchmarks, 2024. [arXiv:2407.20060](#).
5. L. Prokhorenkova, G. Gusev, A. Vorobev, A. V. Dorogush, and A. Gulin. **CatBoost: Unbiased Boosting with Categorical Features.** NeurIPS, 2018. [arXiv:1706.09516](#).
6. M. Raasveldt and H. Mühleisen. **DuckDB: An Embeddable Analytical Database.** SIGMOD, 2019. [doi:10.1145/3299869.3320212](#).

7. N. Hollmann, S. Müller, L. Purucker, A. Krishnakumar, M. Körfer, S. B. Hoo, R. T. Schirrmeister, and F. Hutter. **Accurate Predictions on Small Data with a Tabular Foundation Model.** *Nature* 637, 319–326, 2025. [doi:10.1038/s41586-024-08328-6](https://doi.org/10.1038/s41586-024-08328-6).
8. Prior Labs Team. **TabPFN-3: Technical Report.** 2026. [arXiv:2605.13986](https://arxiv.org/abs/2605.13986).
9. Y. Zhang, L. Xu, Q. Gan, D. Wipf, and M. Wang. **RDBLearn: Simple In-Context Prediction over Relational Databases.** 2026. [arXiv:2602.18495](https://arxiv.org/abs/2602.18495).
10. A. Hayler, K. Flöge, A. Arazi, R. Ranjan, J. Leskovec, L. Purucker, F. Hutter, N. Hollmann, and the Prior Labs Team. **Advancing Open and Reproducible Relational Learning: RelArena- α , TabPFN-Rel and RPI.** 2026. [arXiv:2608.16319](https://arxiv.org/abs/2608.16319).
11. W. L. Hamilton, R. Ying, and J. Leskovec. **Inductive Representation Learning on Large Graphs.** NeurIPS, 2017. [arXiv:1706.02216](https://arxiv.org/abs/1706.02216).
12. T. Chen, C. Kanatsoulis, and J. Leskovec. **RelGNN: Composite Message Passing for Relational Deep Learning.** 2025. [arXiv:2502.06784](https://arxiv.org/abs/2502.06784).
13. V. P. Dwivedi, S. Jaladi, Y. Shen, F. López, C. I. Kanatsoulis, R. Puri, M. Fey, and J. Leskovec. **Relational Graph Transformer.** 2026. [arXiv:2505.10960](https://arxiv.org/abs/2505.10960).
14. A. A. Hagberg, D. A. Schult, and P. J. Swart. **Exploring Network Structure, Dynamics, and Function using NetworkX.** Proceedings of the 7th Python in Science Conference, 2008. [conference paper](#).
15. V. Hudovernik, F. López, V. Kocijan, A. Nitta, J. E. Lenssen, J. Leskovec, and M. Fey. **KumoRFM-2: Scaling Foundation Models for Relational Learning.** 2026. [arXiv:2604.12596](https://arxiv.org/abs/2604.12596).

16 Appendix A: Reference API

The generalized interface accepts table sources, keys, dates, relationships, and labels without requiring direct GraphReduce construction:

```
import kurversc

result = kurversc.fit(
    parent_node=parent,
    label_node=labels,
    tables=tables,
    relationships=relationships,
    max_depth=3,
    feature_family_max_columns=4,
    feature_family_max_features_per_column=32,
    feature_propagation_max_functions_per_column=1,
    forward_search_beam_width=2,
    screening_rows=10_000,
    sample_rows=50_000,
    confirmation_top_k=8,
    adaptive_depth_promotion=True,
    capability_pruning=True,
    search_max_features=8_000,
    search_full_data=True,
    infer_ts_periods=False,
    full_training_frames=1,
    rerank_top_k=3,
    rerank_cutoff_frames=3,
    auto_text_features=False,
    model_backend="catboost",
    duckdb_memory_limit="64GB",
    duckdb_max_temp_directory_size="128GB",
)
```



```

print(result.best_config)           # raw metric winner
print(result.recommended_config)    # complexity-aware selected program
print(result.results)               # full search audit
print(result.confirmation_results)  # medium-fidelity finalists
print(result.rerank_results)        # full-frame finalists
print(result.feature_audit)         # source-column funnel audit
predictions = kurversc.predict(
    result,
    parent_node=parent,
    prediction_node=future_entities,
    tables=tables,
    relationships=relationships,
)
RelBench tasks use the official adapter:
problem = kurversc.load_relbench_problem(
    dataset_name="rel-stack",
    task_name="user-badge",
)

result = kurversc.fit(**problem.fit_kwargs())

```

17 Appendix B: Default meta-parameter frontier

Optional-family level	Potential subsets/configurations at narrow cap	Maximum beam-admitted	Maximum promoted at wide cap	
0: base only	1 / 6	4	at most 2 across all levels	
1 added family	4 / 20	8		—
2 added families	6 / 26	6		—
3 added families	4 / 16	4		—
4 added families	1 / 4	2		—
Total	16 / 72	24		2

An explicitly requested second source-column budget contributes another 72 potential audit records. The default remains one fixed budget.

18 Appendix C: Artifact and trial record

A completed candidate trial records:

- GraphConfig and nominal complexity;
- task metric and validation score;
- objective score used for ordering;
- train and validation row counts;
- feature count and ordered feature names;
- categorical and datetime feature roles;
- feature-construction and model-fit seconds;
- frozen execution plan;
- negligible-gain or drastic-growth notes; and
- completion, skip, or failure status with error text.

The fitted artifact additionally records the chosen graph configuration, production plan, plan fingerprint, downstream model or ensemble, target class mapping, validation score, test predictions and score when available, and the schema required to prepare future prediction frames.